

Pipeline stages: What, Why, Who, and How (task sheet Exercise 1 Step 2)

Stage	What (inputs/output)	Why	Who	How (internal steps/tools)
Data ingestion	Input: new .pgn.zst Output: bronze objects + job metadata	Source of truth late binding to schema catch drift early	Data Engineer	1. Apache Airflow 2. Checksum validation 3. zstd decompression 4. Spark parser 5. Land to Minio bucket
Data pre-processing	Input: raw pgn Output: Delta tables, QA report	ML & SQL need clean types testable contracts	Data Engineer, ML Engineer	1. PySpark 2. Great Expectations 3. MinIO 4. Feast
Model development	Input: versioned feature slice Output: model + metrics in MLflow	Reproducible training avoid skew-serve mismatch	ML Engineer. Data Scientist	1. Temporal train/val/test split 2. LightGBM/XGBoost 3. MLflow (tracking + registry) 4. SHAP explainability 5. Registry gate (metric threshold)
Model deployment	Input: registered model Output: prediction service / batch	Deliver business value controlled rollout	ML Engineer MLOps	1. FastAPI 2. Blue-green deployment 3. shadow traffic / canary 4. Automatic rollback
Model monitoring	Input: live/batch scores + labels Output: alerts, retrain jobs	Drift and Data debt are ML-specific fallout models	SRE, ML engineer, data scientist on-call	1. Evidently (drift, data quality) 2. PSI / KS for drift detection 3. Grafana, Prometheus 4. Airflow retrain DAG